

Exploring the Effects of Hardware Heterogeneity on FaaS Performance Variability

Jonas Heisterberg | 3S | Master Thesis | 19. December 2025

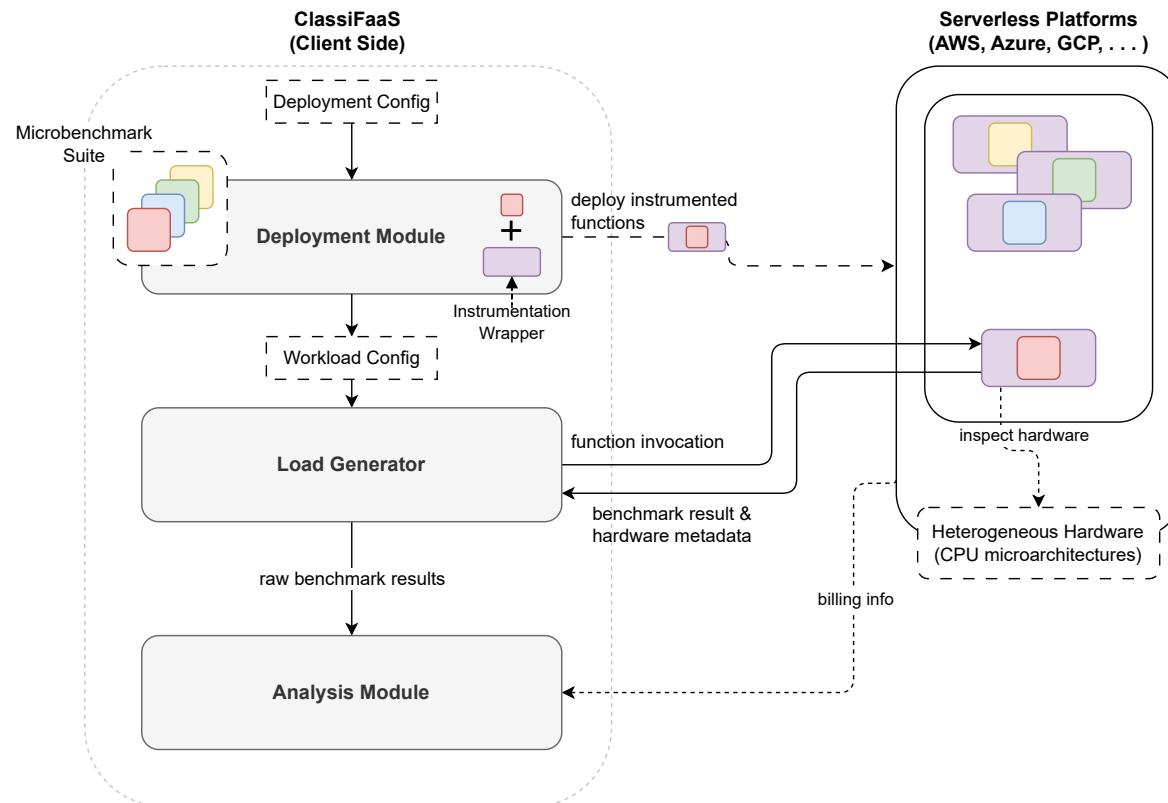
Platzhalter für Sublogo
in Weiß

Motivation: Hardware is abstracted, variability is not



- FaaS abstracts infrastructure from developers
- Same function, same config → different performance
- Pay-per-use billing → performance variance = cost variance
- **Hypothesis:** Identical configurations are backed by heterogeneous CPU generations, causing performance variation
- **Goal:** quantify the “hardware lottery” effect on performance

Benchmark Design



- **Workload:**
 - Closed model
 - Synthetic microbenchmarks
- **Maximize hardware exposure by:**
 - Sequential benchmark execution
 - High concurrency
 - Forced instance termination

CPU Identification Across Platforms



Provider	AWS	Google Cloud Platform	Alibaba Cloud	Azure
Service	Lambda	Cloud Functions (gen 1)	Function Compute 3.0	Functions (Flex Consumption)
CPU identification from <code>/proc/cpuinfo</code>	model name (generic)	numeric model identifier	model name (sometimes generic)	model name (unmodified)
Identifier example	Intel Xeon @ 2.50GHz	85	Intel Xeon 2.50GHz or Intel Xeon Platinum 8269CY 2.50GHz	Intel Xeon Platinum 8370C @ 2.80GHz
Distinct CPU types observed	5	6	5	3

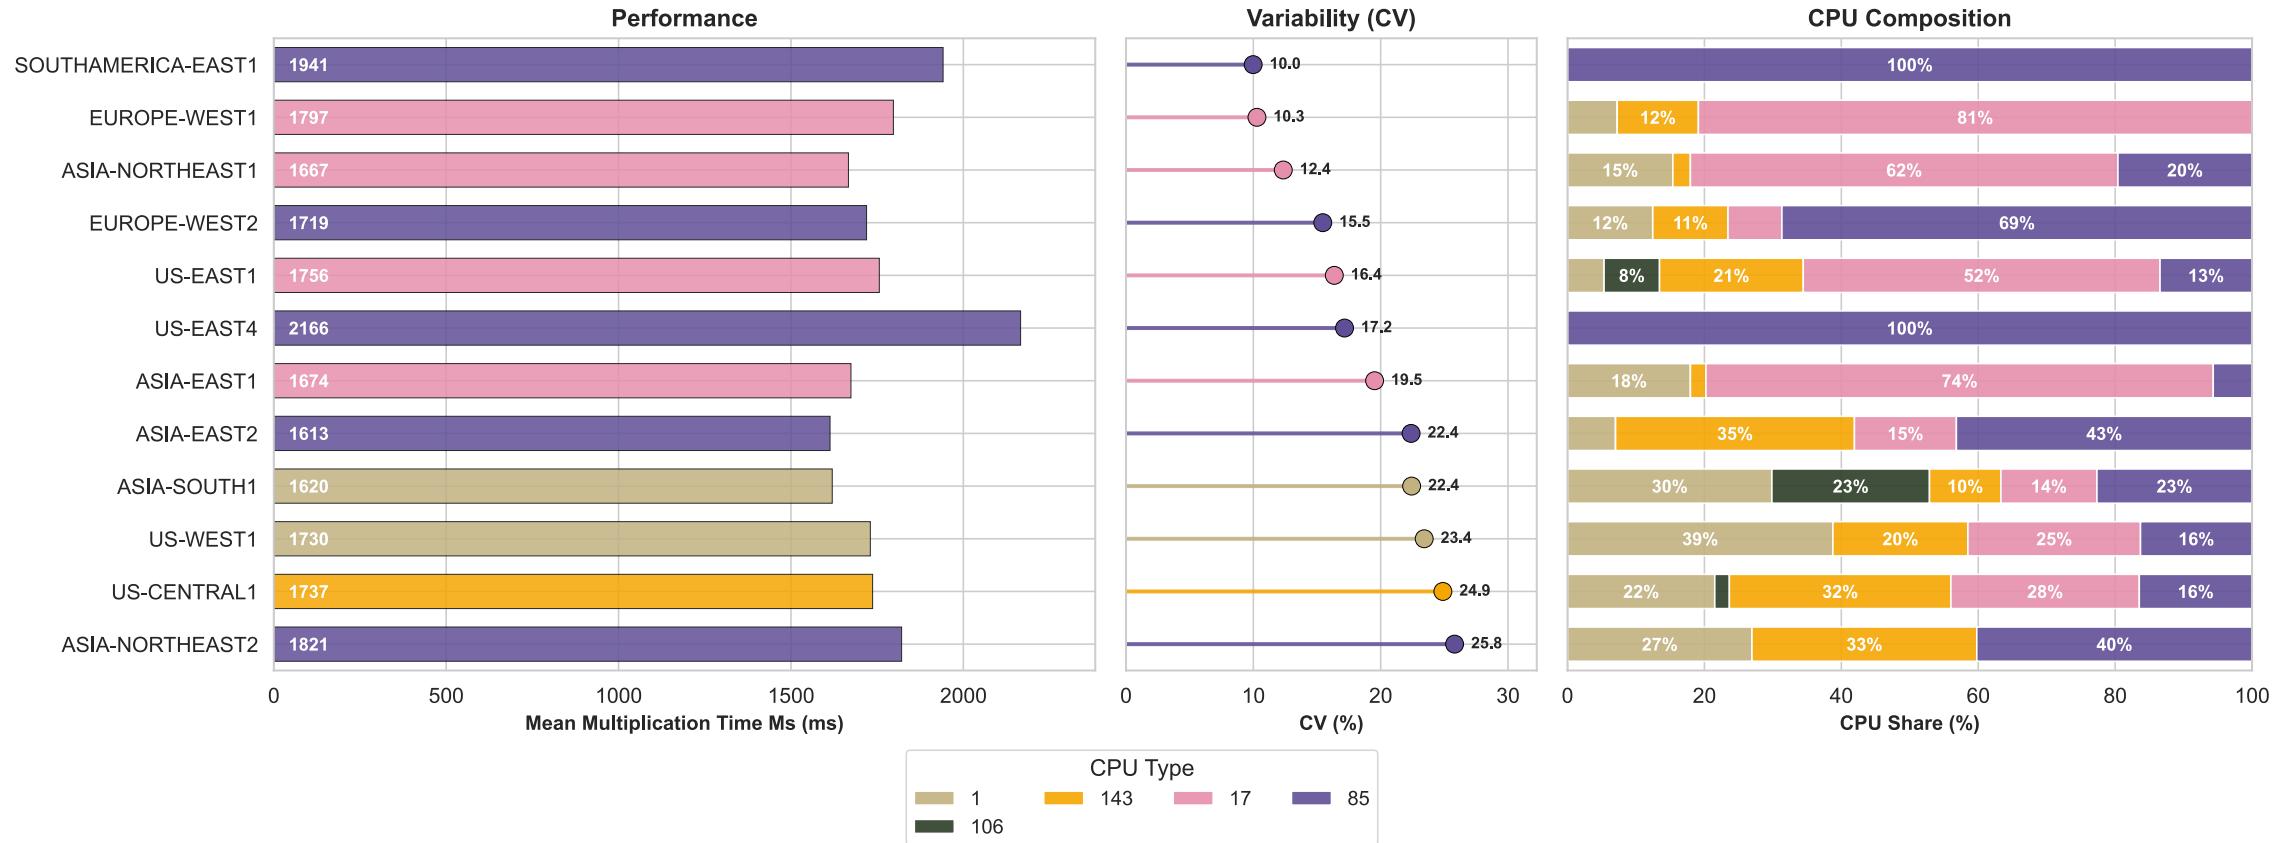
Experimental Design: Progressive Stages



- Stage A — Regional Heterogeneity Baseline
 - Investigate CPU diversity across regions (12 regions per provider)
- Stage B — Memory Configuration Effects
 - Assess the impact of configuration choice (1 region x 3 configurations)
- Stage C — Temporal & Workload Effects
 - Explore hardware assignment over time (1 region x 4 times a day x 7 days)
 - Quantify impact on different workloads (5 distinct workloads)

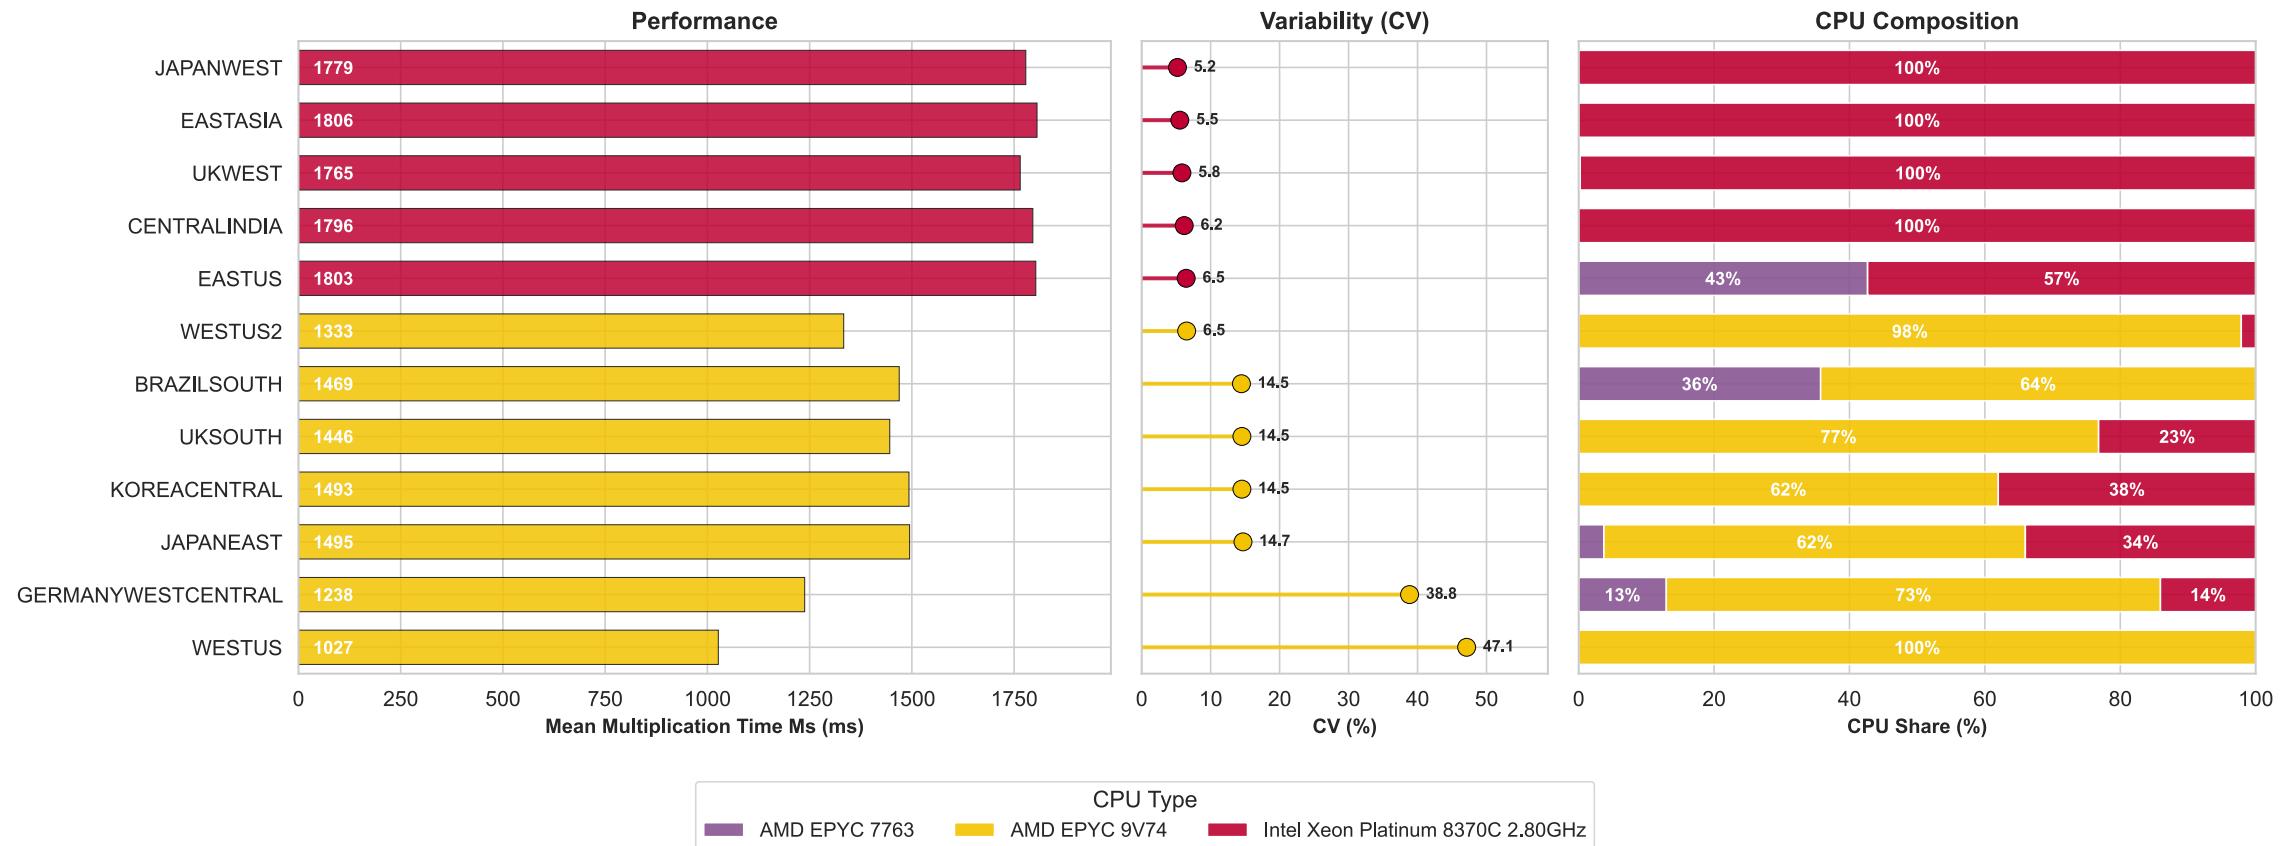
Stage A

GCP - Regional CPU mix and performance variability (Matrix Multiplication, 512 MB)



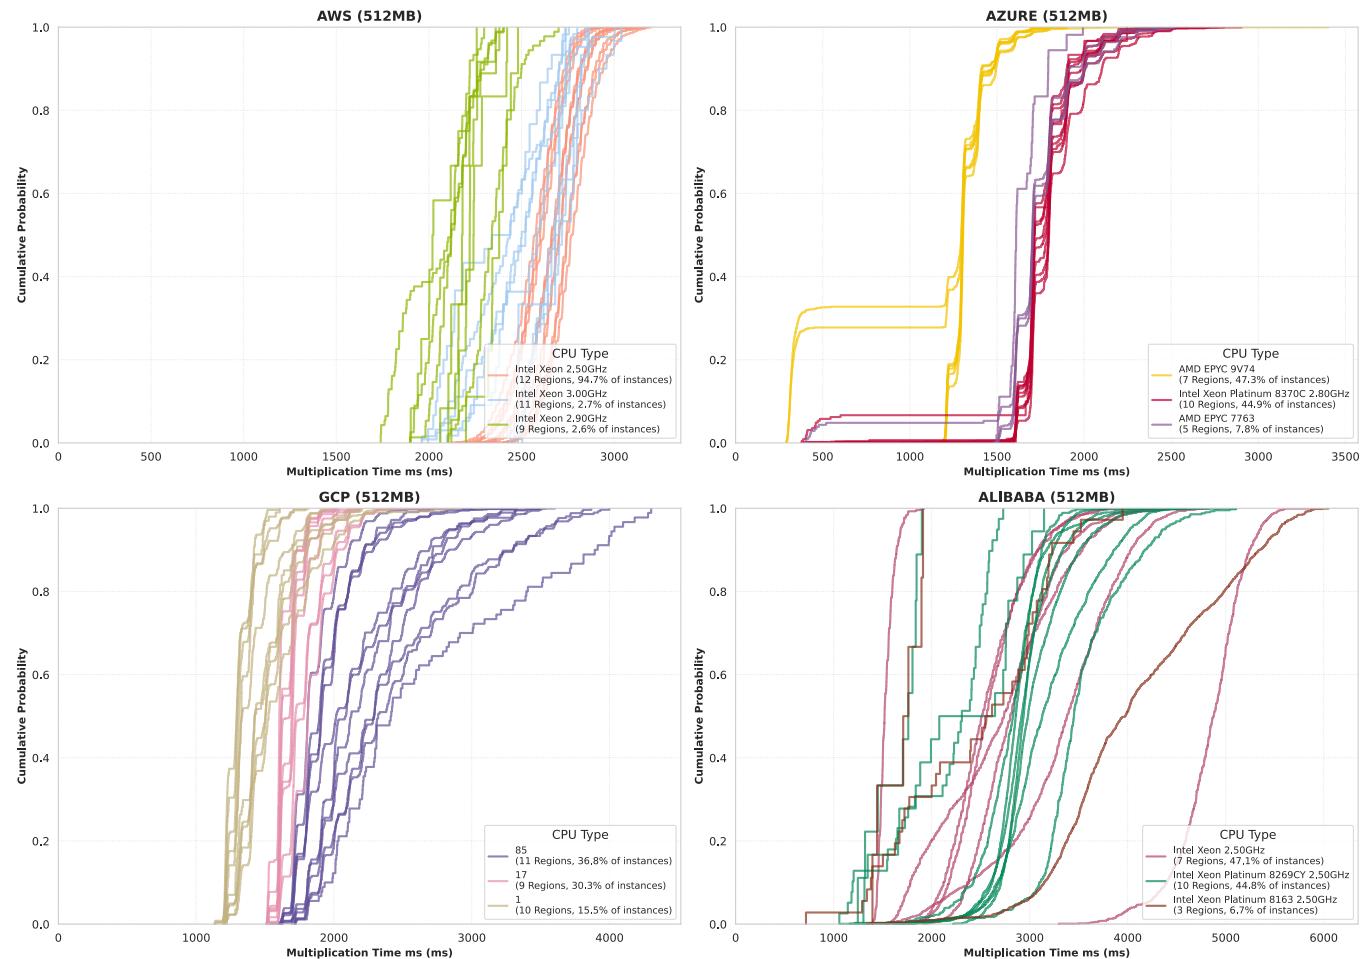
Stage A

AZURE - Regional CPU mix and performance variability (Matrix Multiplication, 512 MB)



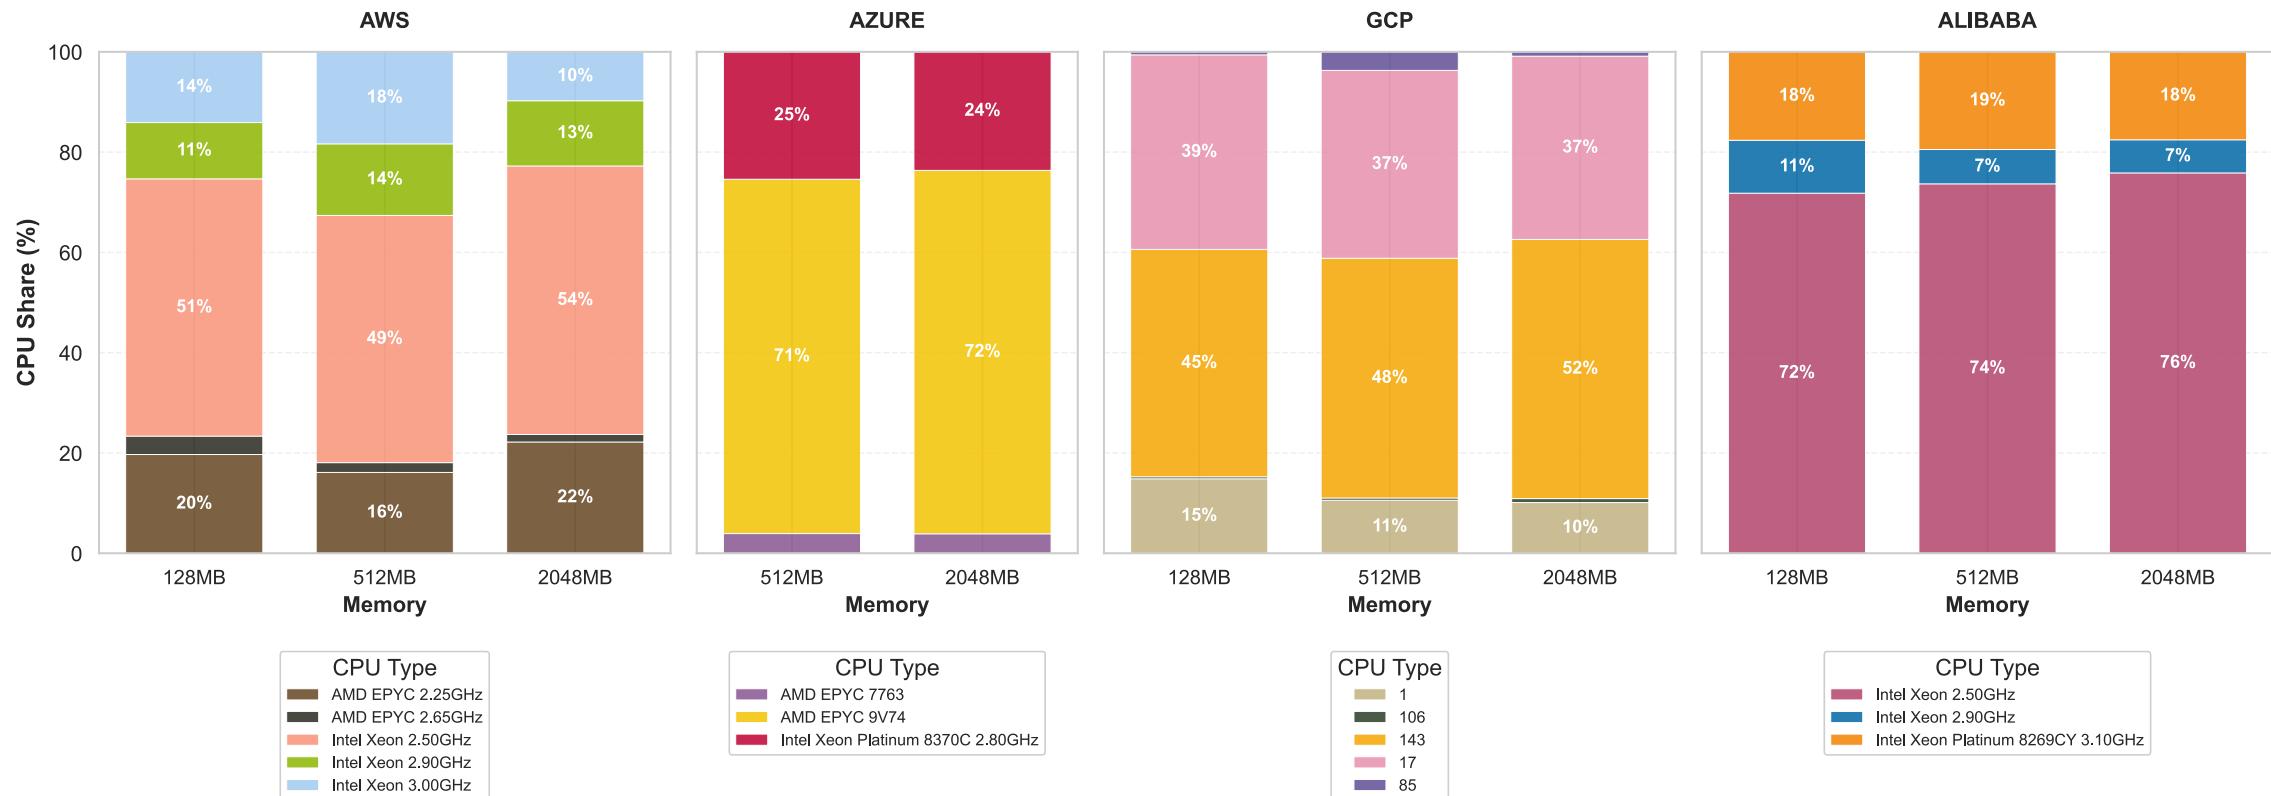
Stage A

ECDF - Regional performance consistency by CPU (Matrix Multiplication)



Stage B

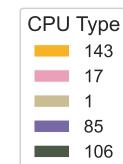
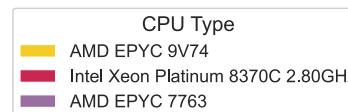
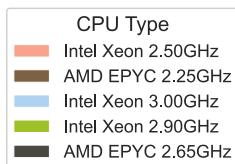
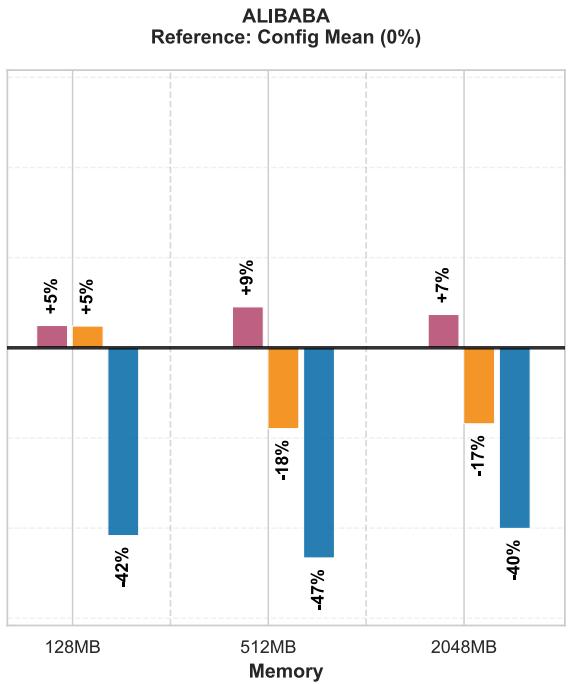
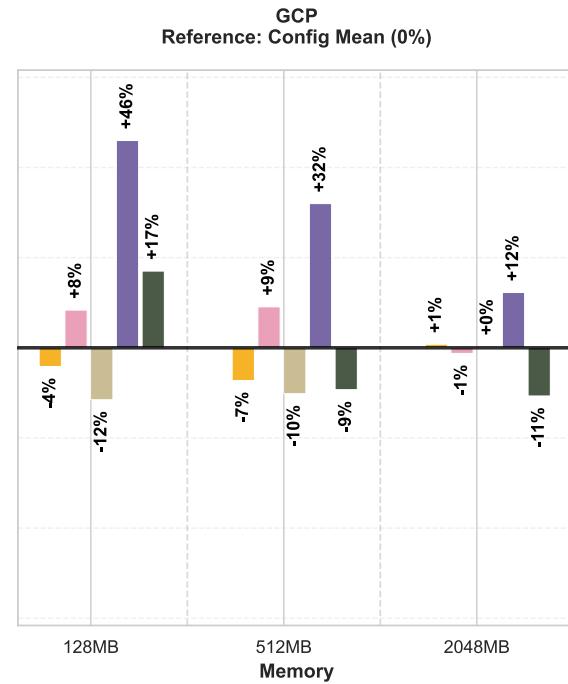
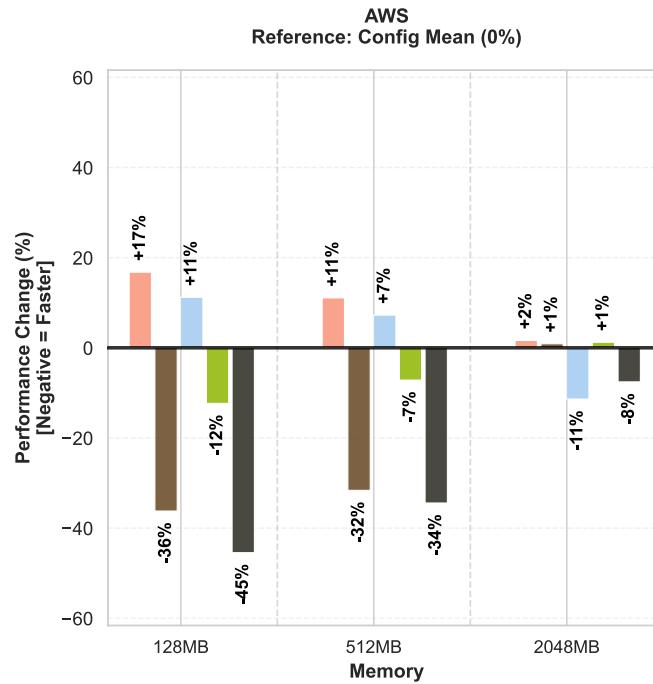
CPU mix by memory size (Matrix Multiplication)



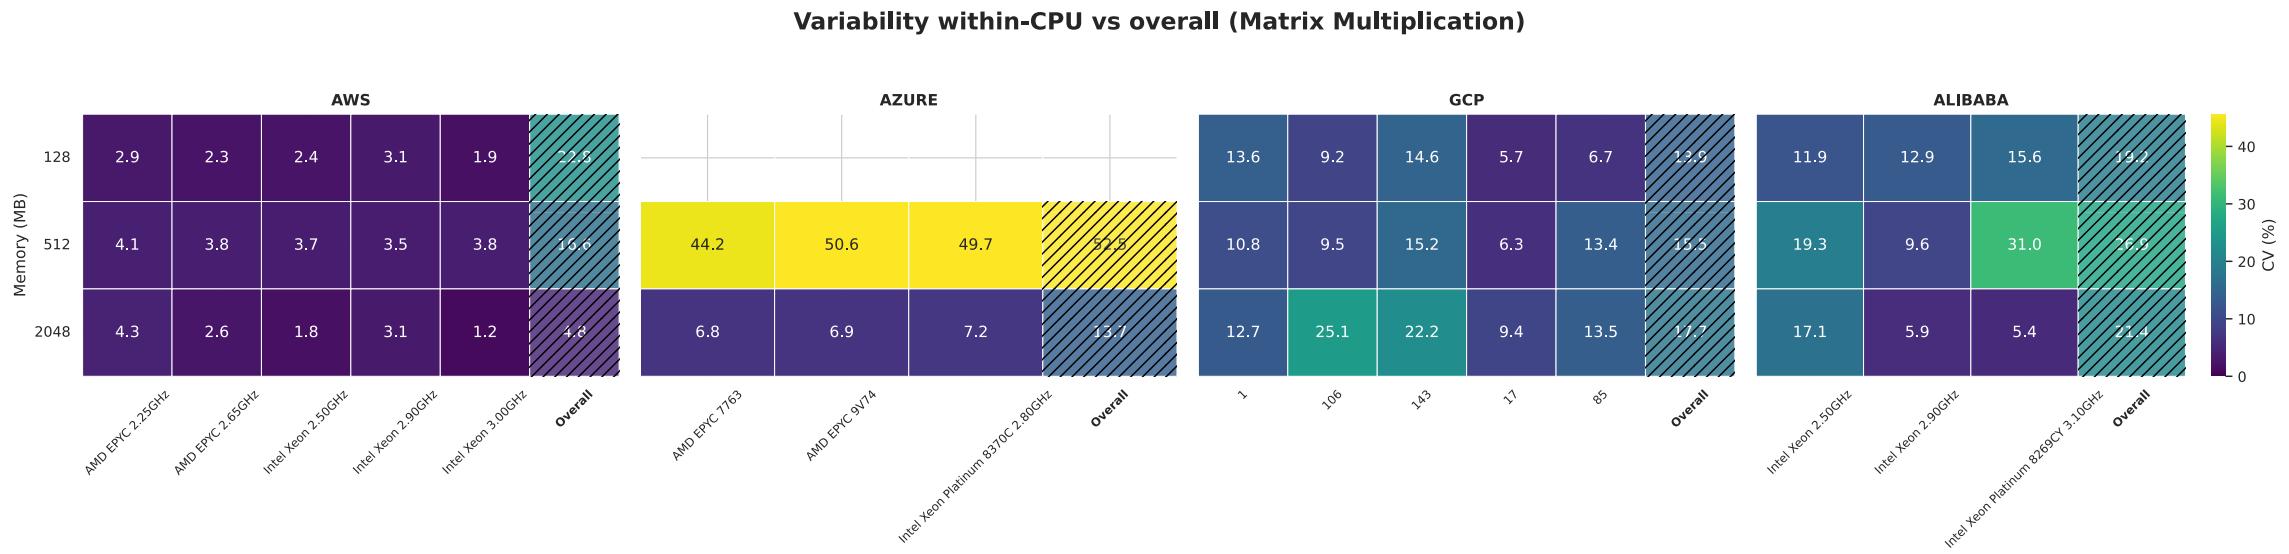
Stage B



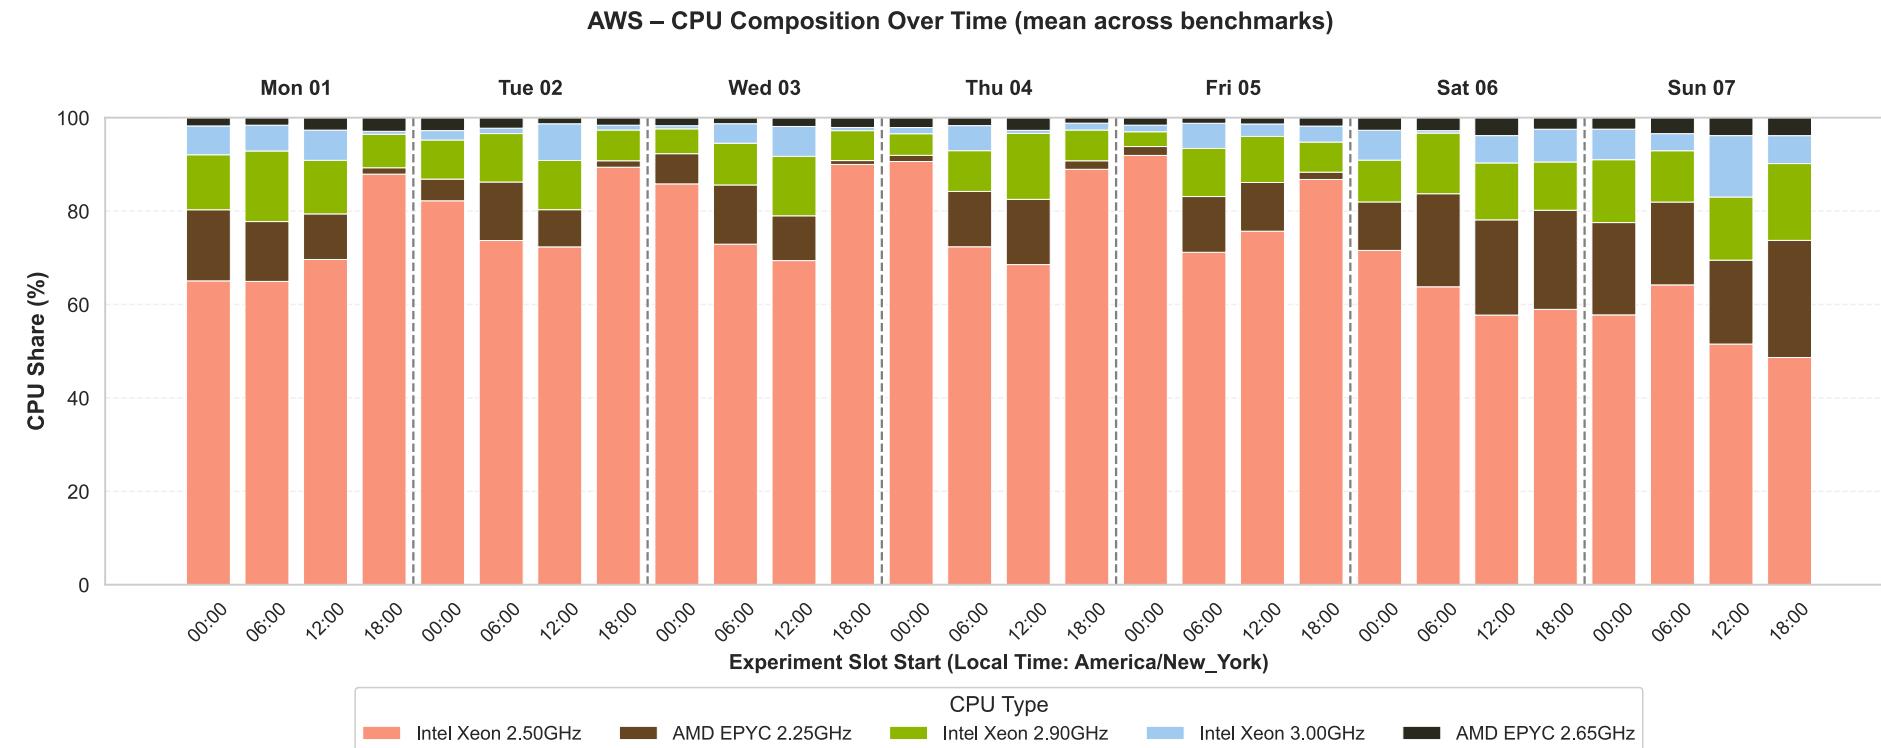
CPU performance deviation from config mean (Matrix Multiplication)



Stage B



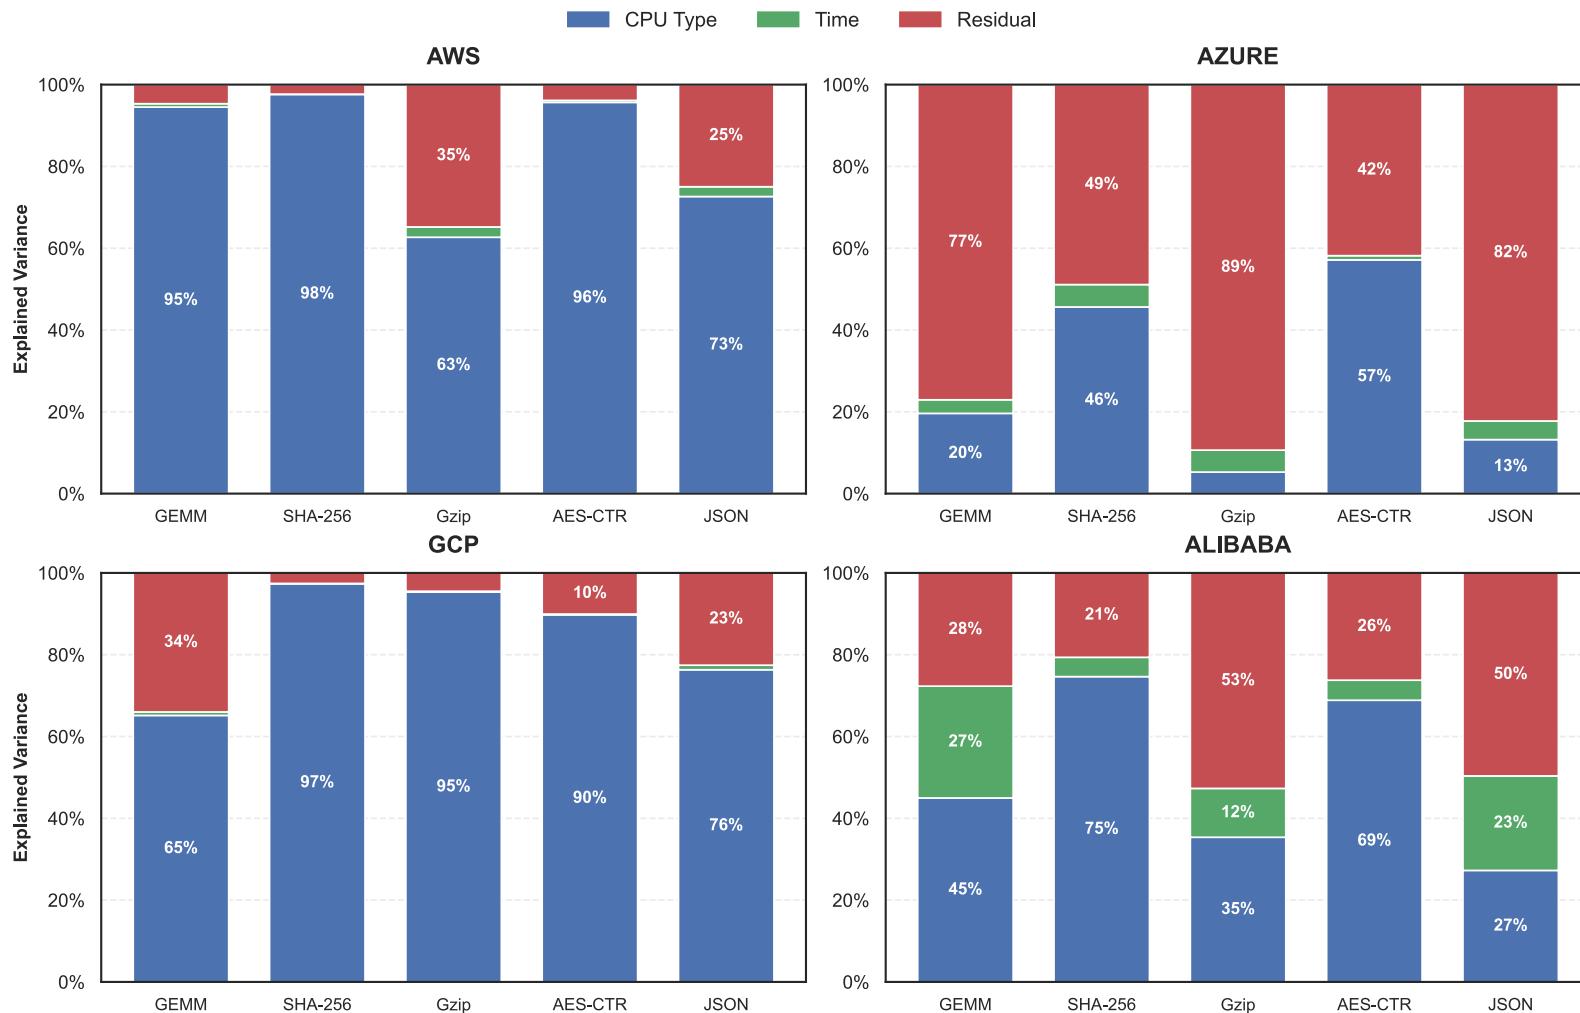
Stage C



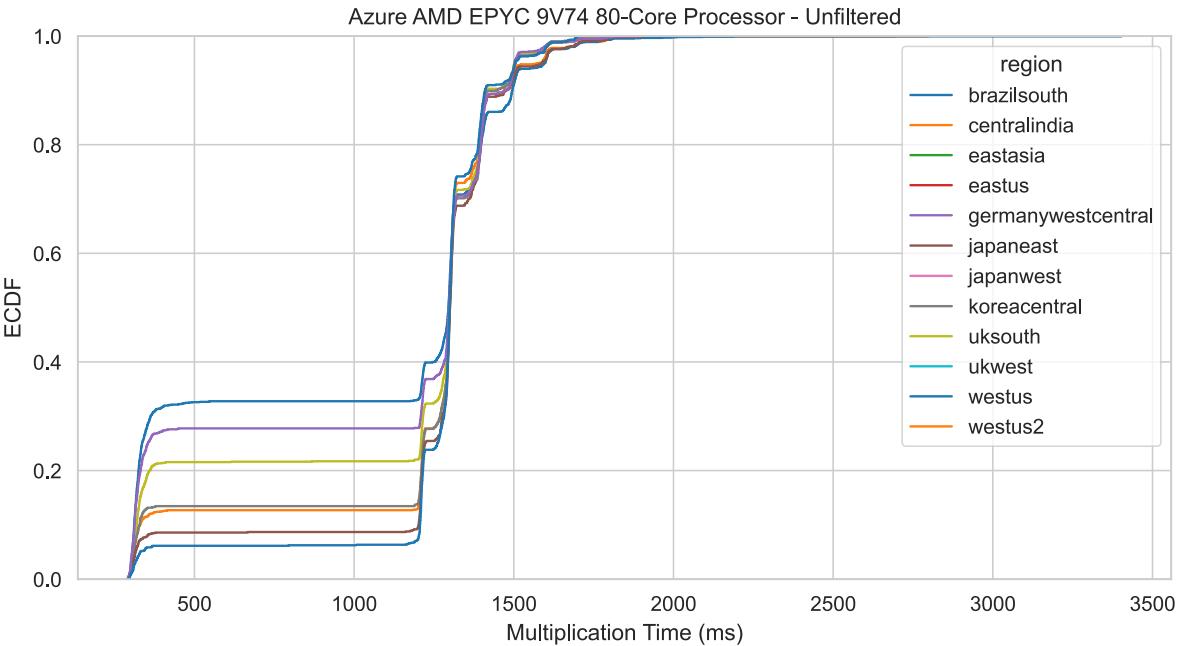
- All Platforms stable except AWS Lambda

Stage C

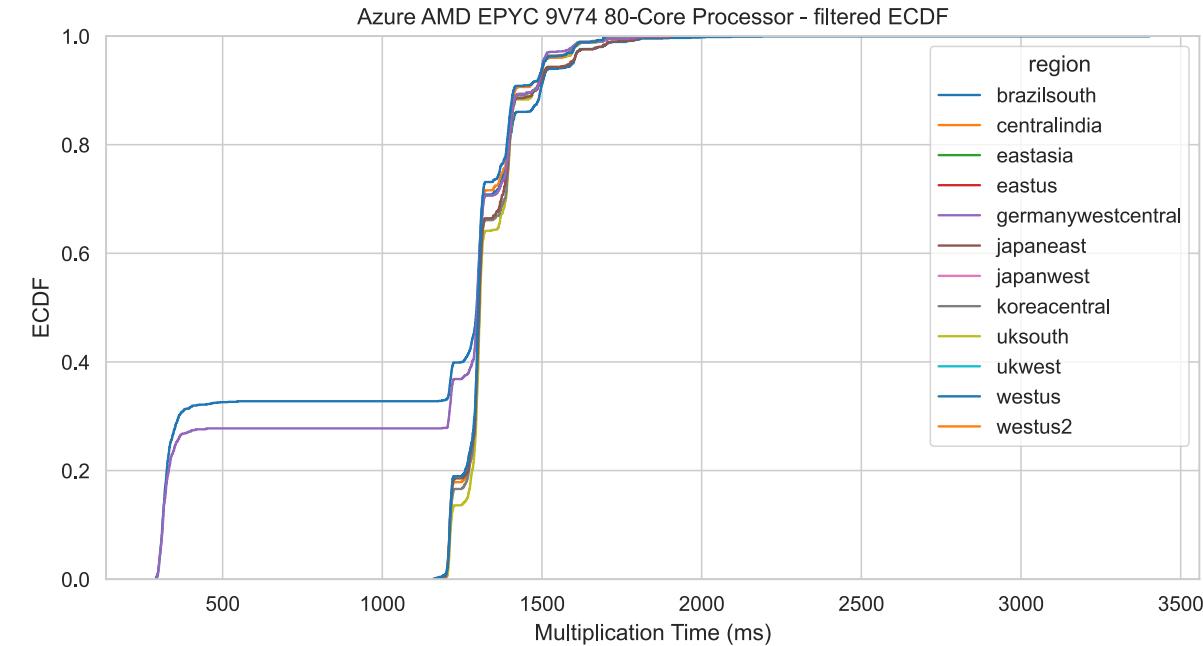
Variance Decomposition: 512MB RAM
(Type-II ANOVA)



Problem: Tukey Outlier Filtering



Raw Data



After Tukey Filter